


Artificial Intelligence: A Modern Approach

Fourth Edition



Chapter 4

Search in Complex Environments

Outline

- ◆ Local Search and Optimization Problems
 - ◆ Hill-climbing
 - ◆ Simulated annealing
 - ◆ Genetic algorithms
- ◆ Local search in continuous spaces
- ◆ Search with Nondeterministic Actions
- ◆ Search in Partially Observable Environments

Local Search and Optimization Problems

In many optimization problems, **path** is irrelevant; the goal state itself is the solution

Then state space = set of “complete” configurations; find **optimal** configuration, e.g., TSP
or, find configuration satisfying constraints, e.g., timetable


In such cases, one can use **iterative improvement** algorithms; keep a single “current” state, try to improve it

Local search algorithms operate by searching from a start state to neighboring states, without keeping track of the paths, nor the set of states that have been reached.

They are not systematic—they might never explore a portion of the search space where a solution actually resides.

Example: Travelling Salesperson Problem

Start with any complete tour, perform pairwise exchanges




Variants of this approach get within 1% of optimal very quickly with thousands of cities

Example: n -queens

Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

Move a queen to reduce number of conflicts



Almost always solves n -queens problems almost instantaneously for very large n , e.g., $n = 1\text{million}$

Hill-climbing (or gradient ascent/descent)


“Like climbing Everest in thick fog with amnesia”

```
function Hill-Climbing( problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← Make-Node(Initial-State[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if Value[neighbor] ≤ Value[current] then return State[current]
    current ← neighbor
  end
```

Hill-climbing contd.

Useful to consider state space landscape



Random-restart hill climbing overcomes local maxima—trivially complete



Random sideways moves escape from shoulders



loop on flat maxima

Simulated annealing

Idea: escape local maxima by allowing some “bad” moves
but gradually decrease their size and frequency

```
function Simulated-Annealing( problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”
  local variables: current, a node
                   next, a node
                   T, a “temperature” controlling prob. of downward steps

  current ← Make-Node(Initial-State[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{Value}[\textit{next}] - \text{Value}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```


Properties of simulated annealing

At fixed “temperature” T , state occupation probability reaches Boltzman distribution

$$p(x) = \frac{e^{-\frac{E(x)}{kT}}}{\sum_y e^{-\frac{E(y)}{kT}}}$$

T decreased slowly enough \Rightarrow always reach best state x^*
 because $\frac{e^{-\frac{E(x^*)}{kT}}}{e^{-\frac{E(x)}{kT}}} = e^{\frac{E(x) - E(x^*)}{kT}} \approx 1$ for small T

Is this necessarily an interesting guarantee??

Devised by Metropolis et al., 1953, for physical process

modelling Widely used in VLSI layout, airline scheduling, etc.

Local beam search

Idea: keep k states instead of 1; choose top k of all their successors

Not the same as k searches run in parallel!

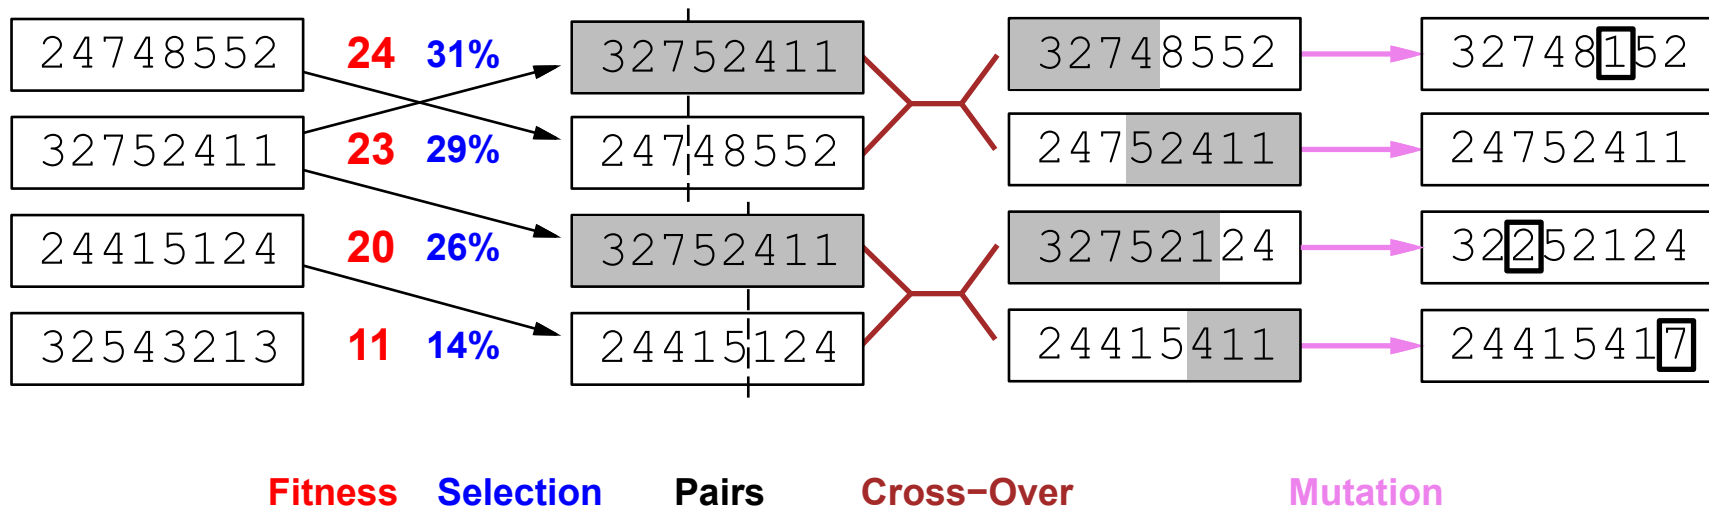
Searches that find good states recruit other searches to join them

Problem: quite often, all k states end up on same local hill

Idea: choose k successors randomly, biased towards good ones Observe the close analogy to natural selection!

Genetic algorithms


= stochastic local beam search + generate successors from **pairs** of states



Genetic algorithms contd.

GAs require states encoded as strings (GPs use programs)

Crossover helps iff substrings are meaningful components



GAs \neq evolution: e.g., real genes encode replication machinery!

Continuous state spaces

Suppose we want to site three airports in Romania:

- 6-D state space defined by $(x_1, y_1), (x_2, y_2), (x_3, y_3)$
- objective function $f(x_1, y_1, x_2, y_2, x_3, y_3) =$
sum of squared distances from each city to nearest airport

Discretization methods turn continuous space into discrete space, e.g., empirical gradient considers $\pm\delta$ change in each coordinate

Gradient methods compute $\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x_1} & \frac{\partial f}{\partial y_1} & \frac{\partial f}{\partial x_2} & \frac{\partial f}{\partial y_2} & \frac{\partial f}{\partial x_3} & \frac{\partial f}{\partial y_3} \end{bmatrix}$

to increase/reduce f , e.g., by $\mathbf{x} \leftarrow \mathbf{x} + a \nabla f(\mathbf{x})$

Sometimes can solve for $\nabla f(\mathbf{x}) = \mathbf{0}$ exactly (e.g., with one city). Newton–Raphson (1664, 1690) iterates $\mathbf{x} \leftarrow \mathbf{x} - \mathbf{H}_f^{-1}(\mathbf{x}) \nabla f(\mathbf{x})$ to solve $\nabla f(\mathbf{x}) = \mathbf{0}$, where $H_{ij} = \partial^2 f / \partial x_i \partial x_j$

Search with Nondeterministic Actions

Agent doesn't know the state its transitioned to **after action**, the environment is **nondeterministic**.

Rather, it will know the possible states it will be in, which is called "**belief state**"

Examples:

- The **erratic vacuum world** (if-then-else) steps. If statement tests to know the current state.
- **AND-OR** search trees. Two possible actions (**OR nodes**). Branching that happens from a choice (**AND nodes**).
- **Try, try again**. A cyclic plan where minimum condition (every leaf = goal state & reachable from other points in the plan)

Search in Partially Observable Environments

Problem of partial observability, where the agent's percepts are not enough to pin down the exact state.

Searching with no observation: Agent's percepts provide **no information at all**, sensorless problem (or a conformant problem).

- Solution: sequence of actions, not a conditional plan

Searching in partially observable environments requires a function that **monitors** or **estimates** the environment to maintain the belief state.

Summary

Local search methods keep only a **small number of states** in memory that are useful for optimization.

In **nondeterministic environments**, agents can apply **AND–OR search** to generate contingency plans that reach the goal regardless of which outcomes occur during execution.

Belief-state is the set of **possible states** that the agent is in for **partially observable environments**.

Standard search algorithms can be **applied directly to belief-state** space to solve **sensorless** problems.